

ScholAR: Page-Grounded Retrieval-Augmented Generation for Scientific Document Comprehension

Anonymous submission

Abstract

Understanding a scientific paper means being able to trace each claim back to the exact place it came from. General-purpose chat over an uploaded PDF is fluent but unreliable for this: nothing ties a cited page number to the text the model was actually shown, and the scientific retrieval-augmented systems that do better tend to rely on cloud-hosted frontier models or datastores of millions of papers. We present ScholAR (Smart Companion for Holistic Organization, Literature Analysis & Research), a retrieval-augmented generation system for scientific documents that runs entirely on a consumer laptop through a local model server, with no cloud API for text or vision. Two design choices carry the system. Page-preserving chunking keeps every retrieved unit inside a single PDF page, so a citation resolves to one highlightable location. Indirect citation grounding lets the model cite only evidence identifiers that map to chunks it actually retrieved, so it cannot write a page number of its own. We evaluate retrieval, answer faithfulness, visual grounding, and cross-document localization on labeled benchmarks. On 51 cases, the generated answers stay grounded in their retrieved context with a mean score of 0.971, no answer atom contradicts its source, and 94% of inline citations are fully supported with none unsupported. BM25-primary retrieval reaches Recall@5 of 0.929, and figure and table questions route to the correct visual element in all 18 pilot cases. Cross-document localization sits at or below a random-guess floor, which we report as our clearest open problem. We also contribute a 100-case benchmark across 25 papers spanning text, mathematical, and figure questions, and a model-agnostic human-evaluation protocol over four local models whose expert scoring is under way.

1 Introduction

Reading an unfamiliar paper is slowest at the moments that matter most: checking that a stated number, a method detail, or a figure actually says what you think it says. Uploading the PDF to a general-purpose chat model helps with fluency but brings a new failure mode. Such systems summarize confidently while inventing page references and attributing quotes to places they do not appear, because a generic chat pipeline puts no hard constraint between a citation and the text the model was conditioned on.

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

Purpose-built scientific retrieval-augmented generation (RAG) systems close part of this gap, but they assume resources most readers do not have on their own machine. OpenScholar (Asai et al. 2024), PaperQA2 (Skarlinski et al. 2024), and SciRAG (Ding et al. 2025) depend on cloud-hosted frontier language models, and OpenScholar additionally on a datastore of tens of millions of papers. We ask a narrower question: what does grounded, citable question answering over a scientific document look like under a hard local-only constraint, with a single sub-15-billion-parameter model, no cloud API, and a retrieval design simple enough to audit end to end?

Two ideas carry ScholAR. Page-preserving chunking segments a document so that every retrieved chunk lies entirely within one PDF page, which makes a citation resolve to an exact, highlightable location rather than an approximate span. Indirect citation grounding assigns each retrieved passage a short-lived evidence identifier and instructs the model to cite only those identifiers, never a free-form page number; the application then rewrites each identifier into a numbered reference that maps deterministically to a page. The model decides only which evidence supports a claim, and the evidence-to-page mapping is application logic rather than a model guess. Beyond single-document text, ScholAR routes figure and table questions to a locally run, natively multimodal model and extends to a paper’s own cited references for cross-document questions.

Our contributions are:

- A fully local RAG system for scientific documents that runs both text and vision inference on-device, using page-preserving chunking and an indirect citation mechanism that rules out page-number hallucination by construction.
- An evaluation that scores the faithfulness of the system’s *generated answers* rather than of pre-written gold claims, paired with an automated per-citation support metric. On 51 cases, generated answers reach a mean grounding score of 0.971 with a zero contradiction rate and 94% fully supported citations.
- Automated benchmarks and results for retrieval, retrieval-support faithfulness, visual grounding, and cross-document localization, each reported with its limitations, including a cross-document result at or below a random-guess floor.

- A 100-case benchmark across 25 papers spanning text, mathematical, and figure questions, with gold answers verified against their source passages, and a model-agnostic human-evaluation protocol over four local models. The expert scoring is under way, and its result tables are marked as placeholders here.

2 Related Work

Retrieval-augmented generation. The RAG framework couples a retriever with a generator so that outputs are conditioned on retrieved evidence (Lewis et al. 2020). Lexical retrieval via BM25 (Robertson and Zaragoza 2009) remains a strong baseline, and the BEIR benchmark (Thakur et al. 2021) documents that zero-shot dense retrievers often fail to beat it. ScholAR uses BM25 as its primary retrieval signal and dense retrieval as a fusion component.

Scientific-literature RAG systems. PaperQA and PaperQA2 (Lála et al. 2023; Skarlinski et al. 2024) are agentic systems that report high accuracy on scientific literature question answering using frontier model backends. OpenScholar (Asai et al. 2024) synthesizes literature over a large open datastore, and SciRAG (Ding et al. 2025) adds citation-graph-aware, outline-guided synthesis. All three assume cloud-hosted or large self-hosted models. ScholAR differs by targeting a single local sub-15-billion-parameter model on consumer hardware, trading raw capability for zero data exposure and no API dependency.

Small models that cite their sources. Closest in spirit to our citation mechanism, Pleias-RAG (Pleias 2025) trains small (350M and 1.2B) models that emit literal source quotes during generation. That line shows citation behavior can be learned into a compact model. ScholAR takes a complementary, model-agnostic route: the grounding lives in the retrieval and application layer through evidence identifiers, so it works with any local generation model rather than requiring a model trained to quote.

Citation and faithfulness evaluation. Gao et al. (2023) formalize citation precision and recall for generated text using entailment between cited passages and statements. SummaC (Laban et al. 2022) scores summary consistency with a zero-shot, sentence-level entailment variant that needs no fine-tuned model, and FActScore (Min et al. 2023) decomposes text into atomic facts verified against a source. RAGAS (Es et al. 2024) automates RAG faithfulness scoring against retrieved context. Our faithfulness metric adopts the SummaC-style zero-shot entailment and the FActScore-style atomic decomposition, and we apply it to the system’s own generated answers.

Visual and multi-document scientific QA. ColPali (Faysse et al. 2024) indexes pages as visual embeddings for retrieval without OCR. ScholAR keeps text retrieval primary and routes only figure and table questions to a cropped-region vision call, which keeps the common text path fast and fully lexical. M3SciQA (Li et al. 2024) pairs each anchor paper with its cited references and finds that even large foundation models underperform human experts

at cross-paper reasoning, a difficulty our cross-document results echo. Document parsing tools such as GROBID (Lopez 2009) and Nougat (Blecher et al. 2023) target structured extraction and are complementary to our chunking.

3 Problem Formulation

Given a scientific PDF D , a user question q , and a set of extracted chunks $\mathcal{C} = \{c_1, \dots, c_n\}$ derived from D (and, in multi-document mode, from D ’s cited references), the system retrieves a ranked subset $\mathcal{C}_q \subseteq \mathcal{C}$ and generates an answer $\hat{a}$ grounded in that evidence:

$$\hat{a} = \arg \max_a P(a \mid q, \mathcal{C}_q), \quad \mathcal{C}_q = \text{Retrieve}(q, \mathcal{C}, k). \quad (1)$$

Each chunk c_i carries a source page number and, in multi-document mode, a source paper identifier, so $\mathcal{C}_q$ always resolves to a specific page of a specific document. Generation attaches each claim in $\hat{a}$ to an evidence identifier referencing a specific $c_i \in \mathcal{C}_q$ rather than a free-form page number, so citations are constrained to what was actually retrieved, independent of Retrieve’s own precision.

4 Method

Figure 1 shows the end-to-end pipeline. A PDF is chunked one page at a time, ranked by BM25, and exposed to the model as evidence identifiers; the model’s identifier citations are then rewritten into page-anchored references by the application rather than by the model.

4.1 Document Processing and Page-Preserving Chunking

Given an arXiv identifier or an uploaded PDF, ScholAR extracts per-page text with PyMuPDF (Artifex Software 2024) and persists metadata, pages, chunks, and figures as JSON. Standard character or token windows routinely split a chunk across a page boundary, breaking the one-to-one map between a retrieved chunk and the page a citation should highlight. ScholAR instead chunks strictly within each page: it slides a window of target size 1400 words with 120-word overlap over each page’s tokens independently, so a chunk never spans two pages. A lightweight regular-expression pass tags each chunk with a detected section title when a heading is found near the top of the page. Chunk identifiers are numbered per paper.

4.2 BM25-Primary Retrieval

The production retriever ranks chunks with Okapi BM25 (Robertson and Zaragoza 2009) ($k_1=1.4$, $b=0.72$) as the primary signal, then applies small additive heuristic boosts: a query-expansion dictionary maps research terms to related vocabulary, an explicit page number in the query adds a page-hint boost, and query language implying a section or a claim-bearing phrase adds a small boost. A separate hybrid configuration fuses BM25 with dense retrieval over a local all-MiniLM-L6-v2 sentence encoder (Reimers and Gurevych 2019) using Reciprocal Rank Fusion (Cormack, Clarke, and Buettcher 2009) ($k=60$); this configuration is used in the faithfulness study.




Figure 1: The ScholAR pipeline. Page-preserving chunking keeps every retrieved unit inside one PDF page. The model may cite only evidence identifiers $[E_i]$, which the application rewrites into numbered references and maps deterministically to a page; the model never writes a page number, so page-citation hallucination is prevented by construction. Figure and table queries route to the same local multimodal model with a cropped, $3\times$ -zoomed image.

4.3 Indirect Citation Grounding

Asking a model to write a page number directly (for example “[p. 7]”) is fragile, because the model has no reliable way to know which page a claim came from and will produce a plausible-looking one. ScholAR instead assigns each retrieved passage a short-lived evidence identifier for the current turn and instructs the model to cite only these identifiers. After generation, the system discards any page number the model wrote on its own, rewrites each valid evidence identifier into a compact numbered reference in citation order, and resolves each reference to its source chunk’s page for highlighting. The model decides only which evidence supports a claim; the mapping from evidence to page is deterministic application logic. This prevents page-number hallucination by construction rather than by asking the model to be careful.

4.4 Visual Grounding

When a query references visual content, through cue terms such as “figure”, “table”, or an explicit figure number, retrieval boosts figure and table chunks so they can compete against longer text chunks. Each caption is detected by regular expression, and the region above it is rendered at $3\times$ zoom and cached alongside the caption. A query naming an explicit figure number is matched directly against that figure’s label, which matters because a letter-initial tokenizer reduces “Figure 2” and “Figure 14” to the same token once digits are dropped. When the top-ranked chunk is a figure or table, the cached image is sent, with the caption and supporting text, to the same local model used for text generation (both models we use are natively multimodal), rather than to a separate vision model or cloud API. If the image is missing or the call fails, the system degrades to a caption-only extractive answer.

4.5 Multi-Document Extension

ScholAR resolves an anchor paper’s bibliography and lets the user ingest cited works into the same session. For arXiv papers, references are resolved by arXiv identifier; for uploaded PDFs, by a title search over the last pages, matched against numbered and author-year citation styles. Ingesting a reference applies the same page-preserving chunking and tags every chunk with a source paper identifier. At query time, chunks from all loaded papers merge into one retrieval pool, and each citation is labeled with its source paper.

4.6 Faithfulness Metrics

We evaluate faithfulness with a zero-shot, sentence-level entailment approach in the SummaC style (Laban et al. 2022), decomposing text into atomic claims in the FActScore style (Min et al. 2023). For an atom a_i and a chunk C_j we take the maximum cosine similarity between the atom and the chunk’s sentences under the local all-MiniLM-L6-v2 encoder, and label the atom entailment, neutral, or contradiction against thresholds $\tau_e=0.42$ and $\tau_n=0.25$. A Combined Faithfulness Score (CFS) blends three tiers:

$$\text{CFS} = 0.50 \text{ NLI} + 0.30 \text{ SCR} + 0.20 \text{ KFP}, \quad (2)$$

where NLI is the entailment fraction, SCR is whole-claim to best-chunk cosine similarity, and KFP is boundary-matched recall of numbers and technical terms. The thresholds and the $\text{CFS} \geq 0.55$ faithful cutoff were calibrated on a small pilot set; we did not run a systematic sensitivity sweep.

We apply this metric in two ways that answer different questions. *Retrieval-support CFS* scores a pre-written gold claim against the retrieved chunks, measuring whether retrieval surfaces support for a known fact (answerability). *Generation faithfulness* scores the system’s *actual generated answer* against its retrieved context, measuring whether the output stays grounded and avoids contradicting its source. The second is the quantity our hallucination claim rests on, and it is the result we foreground.

4.7 System Architecture

The backend is a local application that stores parsed PDFs as JSON on disk; the frontend highlights cited pages. All inference, both text generation and figure or table question answering, runs locally through a single natively multimodal model, with no cloud API used for inference and no user data leaving the machine. Paper acquisition and reference resolution use public arXiv and Semantic Scholar lookups by identifier only. All experiments in this paper were run on a consumer laptop (Apple Silicon, 18 GB unified memory) under Python 3.12, with all-MiniLM-L6-v2 loaded from a local cache; no discrete or cloud GPU was used, and inference runs on the laptop’s integrated GPU through the local model server’s accelerated backend.

5 Benchmarks

We use two families of evaluation data. The first is a set of small, manually labeled benchmarks over three landmark

Table 1: Retrieval on 14 manually annotated cases across 3 papers.

System	R@1	R@3	R@5	MRR
Keyword Overlap	0.571	0.786	0.929	0.687
BM25-Only	0.714	0.857	0.929	0.788
BM25+Rerank	0.714	0.857	0.929	0.788
Dense-Only	0.786	1.000	1.000	0.881

NLP papers (Attention Is All You Need (Vaswani et al. 2017), Retrieval-Augmented Generation (Lewis et al. 2020), and LLaMA (Touvron et al. 2023)): a 14-case retrieval benchmark with labeled relevant chunks, a 51-case faithfulness benchmark with oracle claims and a supporting chunk, an 18-case visual-grounding pilot, and an 18-case multi-document benchmark of which 10 have an arXiv-resolvable target. These support the automated metrics that need ground-truth chunk labels, and they are small: the labeled benchmarks cover three papers between them.

The second is a 100-case benchmark spanning 25 papers, built to broaden coverage and reduce the memorization advantage a model has on very famous papers. It contains 50 text, 25 mathematical, and 25 figure or table questions. Each case was mined from a real passage or caption by a local model and then kept only if the required answer facts appear, on a word or number boundary, in the source passage; questions from benchmark-dataset papers, reference lists, and worked examples were filtered out. This diverse benchmark is the dataset for the human evaluation (Section 7).

6 Automated Evaluation

6.1 Retrieval

Table 1 reports retrieval on the 14-case benchmark. BM25-primary reaches Recall@5 of 0.929 and MRR of 0.788, improving on the keyword-overlap baseline by 10.1 points of MRR. On this small benchmark the reranking heuristics and page hints leave the aggregate metrics unchanged, so the gain over the keyword baseline comes from BM25 itself. A single-pass dense-only configuration scores highest here (Recall@5 1.000, MRR 0.881), counter to the general BEIR finding (Thakur et al. 2021); we attribute this to the small size and paraphrased queries of the benchmark, and at faithfulness scale it is hybrid retrieval, not dense-only, that wins (Section 6.2).

6.2 Retrieval-Support Faithfulness

Table 2 reports the retrieval-support CFS on the 51-case benchmark, comparing the BM25-primary retriever with the hybrid BM25+Dense+RRF retriever. Hybrid reaches Combined CFS 0.827 versus 0.807, and improves the supporting-chunk hit rate at depth 5 (SCHR@5) from 0.824 to 0.922. Under the strict $CFS \geq 0.55$ cutoff, BM25-primary classifies 48 of 51 cases and hybrid 49 of 51 as faithful. This metric measures whether retrieval surfaces support for a known claim, not whether the generated answer is faithful; the latter is Section 6.3.

Table 2: Retrieval-support faithfulness on 51 oracle-claim cases. NLI is the entailment fraction; SCR is whole-claim cosine; KFP is key-fact recall; SCHR@5 is the supporting-chunk hit rate at depth 5.

Metric	BM25-primary	Hybrid+RRF
NLI (entailment)	0.941	0.961
SCR (cosine)	0.484	0.508
KFP (key-fact recall)	0.956	0.969
Combined CFS	0.807	0.827
SCHR@5	0.824	0.922
Faithful (≥ 0.55)	48/51	49/51

Table 3: Generation faithfulness of ScholAR’s produced answers on 51 cases (generation model gemma4:12b). Citation counts are over 200 inline citations.

Metric	Value
Mean generation faithfulness (atoms entailed)	0.971
Mean contradiction rate	0.000
Citation support rate (supported / checked)	0.94
Citations supported / partial / unsupported	188 / 12 / 0

6.3 Generation Faithfulness

Table 3 reports the faithfulness of the system’s generated answers on the same 51 queries, using the local model gemma4:12b for generation. The mean generation-faithfulness score is 0.971: 97.1% of the atomic claims in the generated answers are entailed by the retrieved context. The mean contradiction rate is 0.0, so no answer atom was classified as contradicting its source. Of 200 inline citations checked, 188 are fully supported, 12 partially supported, and none unsupported, a citation-support rate of 0.94. These scores measure grounding to the retrieved context and the absence of outright contradiction, both of which run naturally high for a well-behaved retrieval-augmented system, so we read them conservatively. The zero unsupported-citation count is automated evidence for the indirect citation mechanism, not a claim of perfect correctness against ground truth; assessing that is the job of the human evaluation (Section 7).

6.4 Visual Grounding

On the 18-case visual pilot, retrieval routes to the correct figure or table in all 18 cases, with no case falling back to the caption-only path. We also compare two local multimodal models on a keyword-overlap answer-quality proxy (not an LLM-judge): gemma4:12b scores 0.731 and qwen3.5:9b scores 0.812. A caption-only ablation, sending the same prompt without the image, scores 0.752 versus 0.731 for full vision under gemma4:12b. We do not read this as evidence that the image is unnecessary: the proxy rewards lexical overlap with the question rather than grounded visual correctness, and many answers are also restated in nearby prose. Separating these effects requires a correctness-based judge, which we leave to the human study.

Table 4: Human-evaluation results (in progress). Values are placeholders pending expert scoring; they are not estimated. Dimensions are mean 1–5.

Model	Rel.	Cov.	Faith.	Use.	Cite-F1
qwen3.5:9b	–	–	–	–	–
gemma4:12b	–	–	–	–	–
llama3.1:8b	–	–	–	–	–
mistral:7b	–	–	–	–	–

6.5 Multi-Document Localization

On the 10 arXiv-resolvable cases, ScholAR recovers the correct cited paper in the top 5 half the time (Recall@5 0.50, MRR 0.183) and never at rank 1 (Recall@1 0.00). A random-guess floor over the same candidate pool (on average 8.0 candidate papers per case; floor Recall@5 0.625, MRR 0.340) puts the system at or below chance on this small benchmark. An oracle that restricts retrieval to the anchor plus the single correct paper still reaches only Recall@1 0.00 and MRR 0.289, which points to a difficulty that is not only cross-paper competition but also within-paper chunk ranking on these questions. This is the clearest open problem in the system, and it echoes M3SciQA’s finding that even large models struggle at cross-paper localization (Li et al. 2024).

7 Human Evaluation (In Progress)

Automated metrics measure component behavior; they do not tell us whether a reader would judge an answer correct and well-cited. We therefore designed a model-agnostic human evaluation, grounded in the practice of comparable systems (Asai et al. 2024; Ding et al. 2025; Skarlinski et al. 2024), over the 100-case diverse benchmark of Section 5. Each question is answered by four local generation models (qwen3.5:9b, gemma4:12b, llama3.1:8b, mistral:7b) running the *same* ScholAR pipeline, so only the generation model changes. Expert evaluators score every answer on four anchored 1–5 dimensions, Relevance, Coverage, Faithfulness, and Usefulness, and grade each inline citation as Supported, Partial, or Unsupported. Answers are shown blind and in randomized order. From the scores we compute per-model means and citation precision, recall, and F1; a Friedman test across the four models to assess whether grounding is model-agnostic; and the correlation between human Faithfulness scores and the automated faithfulness metric.

The expert scoring is in progress. Table 4 lists the reported quantities as placeholders; the values will be filled from the completed scoring and are not estimated here.

We also report, once available, the Friedman-test result on the Faithfulness dimension (to be filled) and the human-to-automated Faithfulness correlation (to be filled), which together test whether ScholAR’s grounding holds across local models and whether the automated metric tracks expert judgment.

8 Discussion and Limitations

ScholAR’s contribution is a fully local integration whose two genuinely novel elements, page-preserving chunking and indirect citation grounding, are simple and auditable, and in which the generated answers are measurably grounded, with no contradictions and no unsupported citations on the studied set. The limits of that evidence set the agenda for the rest of this section.

The labeled automated benchmarks cover three papers and are small (10 to 51 cases), so their point estimates carry no variance or significance testing, and several deltas fall within noise. The 100-case, 25-paper diverse benchmark broadens coverage, but we lean on it only lightly for automated metrics, since its cases are not labeled with ground-truth chunks; its main role is to back the human study. The generation-faithfulness metric rewards grounding to retrieved context and the absence of contradiction, which leaves it insensitive to subtle unfaithfulness, and that gap is exactly what the human evaluation exists to close.

We do not run a head-to-head comparison against OpenScholar, PaperQA2, or SciRAG. Those systems depend on cloud-hosted models and large datastores, so a fair comparison would require either a cloud run, contradicting our local-only design, or crippling them to an untested local backend. We leave a resource-matched external comparison, for example PaperQA2 on a local backend, to future work. Finally, cross-document localization is at or below a random floor and the oracle result suggests a within-paper ranking difficulty on these questions, which is our clearest open problem.

9 Conclusion

We presented ScholAR, a fully local retrieval-augmented system for grounded scientific document comprehension that combines page-preserving chunking with indirect citation grounding and extends to visual and multi-document questions. On 51 cases the system’s generated answers are grounded with a mean score of 0.971, no contradictions, and 94% fully supported citations with none unsupported; retrieval reaches Recall@5 of 0.929; and figure and table routing is correct in all 18 pilot cases. Cross-document localization is at or below chance and is our clearest open problem. We contribute a 100-case, 25-paper text, math, and figure benchmark and a model-agnostic human-evaluation protocol over four local models, whose expert scoring is in progress and will replace the placeholders reported here.

References

- Artifex Software. 2024. PyMuPDF: Python Bindings for MuPDF.
- Asai, A.; He, J.; Shao, R.; Shi, W.; Singh, A.; Chang, J.; Lo, K.; Soldaini, L.; Feldman, S.; D’arcy, M.; et al. 2024. OpenScholar: Synthesizing Scientific Literature with Retrieval-augmented Language Models. *arXiv preprint arXiv:2411.14199*.
- Blecher, L.; Cucurull, G.; Scialom, T.; and Stojnic, R. 2023. Nougat: Neural Optical Understanding for Academic Documents. In *arXiv preprint arXiv:2308.13418*.

- Cormack, G. V.; Clarke, C. L. A.; and Buettcher, S. 2009. Reciprocal Rank Fusion Outperforms Condorcet and Individual Rank Learning Methods. In *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 758–759.
- Ding, H.; Zhao, Y.; Hu, T.; Patwardhan, M.; and Cohan, A. 2025. SciRAG: Adaptive, Citation-Aware, and Outline-Guided Retrieval and Synthesis for Scientific Literature. *arXiv preprint arXiv:2511.14362*.
- Es, S.; James, J.; Anke, L. E.; and Schockaert, S. 2024. RA-GAS: Automated Evaluation of Retrieval Augmented Generation. *arXiv preprint arXiv:2309.15217*.
- Faysse, M.; Sibille, H.; Wu, T.; Viaud, G.; Hudelot, C.; and Colombo, P. 2024. ColPali: Efficient Document Retrieval with Vision Language Models. In *The Thirty-Eighth Annual Conference on Neural Information Processing Systems*.
- Gao, T.; Yen, H.; Yu, J.; and Chen, D. 2023. Enabling Large Language Models to Generate Text with Citations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 6465–6488.
- Laban, P.; Schnabel, T.; Bennett, P. N.; and Hearst, M. A. 2022. SummaC: Re-Visiting NLI-Based Models for Inconsistency Detection in Summarization. In *Transactions of the Association for Computational Linguistics (ACL)*, volume 10, 163–177.
- Lála, J.; O’Donoghue, O.; Shtedritski, A.; Cox, S.; Rodrigues, S. G.; and White, A. D. 2023. PaperQA: Retrieval-Augmented Generative Agent for Scientific Research. *arXiv preprint arXiv:2312.07559*.
- Lewis, P.; Perez, E.; Piktus, A.; Petroni, F.; Karpukhin, V.; Goyal, N.; Küttler, H.; Lewis, M.; Yih, W.-t.; Rocktäschel, T.; Riedel, S.; and Kiela, D. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, 9459–9474.
- Li, C.; Shangquan, Z.; Zhao, Y.; Li, D.; Liu, Y.; and Cohan, A. 2024. M3SciQA: A Multi-Modal Multi-Document Scientific QA Benchmark for Evaluating Foundation Models. In *Findings of the Association for Computational Linguistics: EMNLP 2024*.
- Lopez, P. 2009. GROBID: Combining Automatic Bibliographic Data Recognition and Term Extraction for Scholarship Publications. In *Research and Advanced Technology for Digital Libraries*, 473–474.
- Min, S.; Krishna, K.; Lyu, X.; Lewis, M.; Yih, W.-t.; Koh, P. W.; Iyyer, M.; Zettlemoyer, L.; and Hajishirzi, H. 2023. FActScore: Fine-Grained Atomic Evaluation of Factual Precision in Long Form Text Generation. *arXiv preprint arXiv:2305.14251*.
- Pleias. 2025. Even Small Reasoners Should Quote Their Sources: Introducing the Pleias-RAG Model Family. *arXiv preprint arXiv:2504.18225*.
- Reimers, N.; and Gurevych, I. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 3982–3992.
- Robertson, S.; and Zaragoza, H. 2009. The Probabilistic Relevance Framework: BM25 and Beyond. *Foundations and Trends in Information Retrieval*, 3(4): 333–389.
- Skarlinski, M. D.; Cox, S.; Laurent, J. M.; Braza, J. D.; Hinks, M.; Hammerling, M. J.; Ponnampati, M.; Rodrigues, S. G.; and White, A. D. 2024. Language Agents Achieve Superhuman Synthesis of Scientific Knowledge. *arXiv preprint arXiv:2409.13740*.
- Thakur, N.; Reimers, N.; Rücklé, A.; Srivastava, A.; and Gurevych, I. 2021. BEIR: A Heterogeneous Benchmark for Zero-Shot Evaluation of Information Retrieval Models. In *Thirty-Fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*.
- Touvron, H.; Lavril, T.; Izacard, G.; Martinet, X.; Lachaux, M.-A.; Lacroix, T.; Rozière, B.; Goyal, N.; Hambro, E.; Azhar, F.; et al. 2023. LLaMA: Open and Efficient Foundation Language Models. *arXiv preprint arXiv:2302.13971*.
- Vaswani, A.; Shazeer, N.; Parmar, N.; Uszkoreit, J.; Jones, L.; Gomez, A. N.; Kaiser, L.; and Polosukhin, I. 2017. Attention Is All You Need. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30.